



**UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH**

ASM - Time Series Project

Authors: Dmitriy Chukhray, Julian Fransen

Date: 2025-01-06

Contents


Introduction	2
1. Identification	2
1a. Data transformations	3
1b.	8
2. Estimation	10
2a.	10
3. Validation	11
3a.	11

Introduction

The dataset under analysis consists of monthly data on victims of traffic accidents in Spain, including fatalities, serious injuries, and minor injuries, recorded on urban and interurban roads. In this project we will apply the Box-Jenkins ARIMA methodology to understand the time-series dynamics of these traffic incidents and to make reliable predictions for future trends. Spanning from 1993 to 2019, the dataset captures over two decades of detailed information, offering a unique opportunity to identify trends, seasonal patterns, and underlying factors that influence traffic accidents.

1. Identification


First, we load the raw data set and plot it to get a first impression of the data. Based on this, we will start with the identification of transformations.




1a. Data transformations

The first step in time series analysis is to check whether transformations are necessary to stabilize variance or remove trends. We start by checking if the data shows constant variance over time. This is important because ARIMA models assume that the variance is constant. To do this, we compute the mean and variance for each year and examine their relationship. Additionally, we visualize yearly boxplots to confirm constant variance over time.

Check of constant variance (variance ~ mean)




Check of constant variance (yearly boxplots)




These plots show correct behaviour: v is basically uncorrelated with m , and the boxplots are similarly sized, implying constant variance. This means that a log transform is not necessary in our case. The next step is to check the existence of a seasonal pattern in the time series. To do that we are using the function `monthplot`.

Seasonal Patterns in Traffic Accident Data



In the plot above, we can clearly observe a seasonal pattern. If there were no seasonal component, the monthly means would remain at approximately the same level over time, and the shapes of the patterns for each month would not exhibit systematic repetition. These periodic fluctuations suggest that certain months consistently experience higher or lower values, driven by underlying seasonal component(s). To account for this seasonality, we apply a seasonal differencing transformation with a yearly period (12 months).

After seasonality transformation



The last step of achieving the stationarity of the time series is to check whether the mean is constant or not. This can be done by examining the plot of the current time series data or by straight forwardly applying regular difference and then examining the change of the time series' variance.

To verify if taking one regular difference is optimal, we calculate the variance of different transformation of the data:

1. original data (`serie`)
2. transformed with yearly season transformation (`d12serie`)
3. one regular difference applied to the previous transformation (`d1d12serie`)
4. another regular difference applied to the previous transformation

The values are here below:

```
## Variance of original data: 1.863327
```

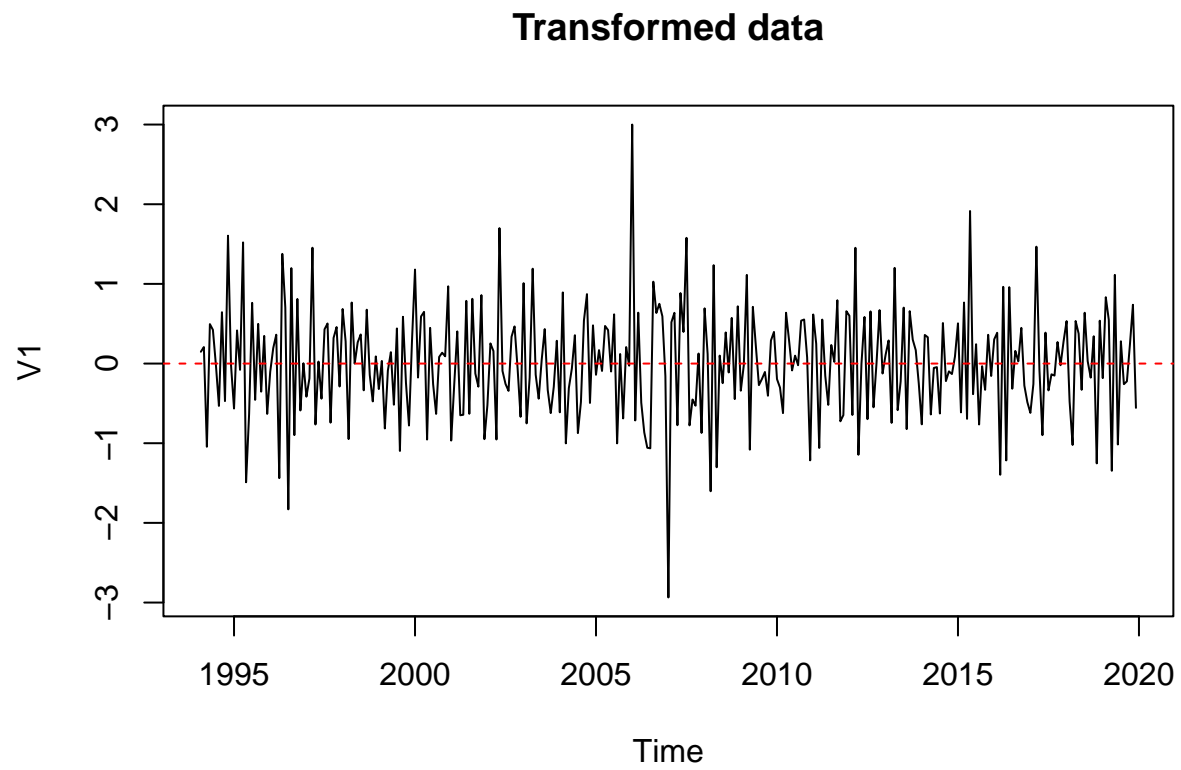
```
## Variance after seasonal differencing: 0.6081206
```

```
## Variance after one regular differencing: 0.4841846
```

```
## Variance after two regular differencing: 1.349775
```


The total variance is minimal for one regular and one 12 month seasonal difference. This means that we should take one regular difference and no more. As we can see below, the data after transformations resembles white noise, which is the aim of data transformations in time-series data analysis. Now that we

have stationary time series, we can propose ARIMA models by evaluating - among other metrics - the ACF and PACF plots.



1b.

Now we are going to take a closer look at the ACF and the PACF plots.



To propose SARIMA models we need to evaluate ACF and PACF plots separately. SARIMA model can be expressed as:

$$SARIMA(p, d, q)(P, D, Q)_s$$

Where:

- p, d, q : Non-seasonal AR, differencing, and MA terms, respectively.
- P, D, Q : Seasonal AR, differencing, and MA terms, respectively.
- s : Periodicity of the seasonal component.

Looking at the ACF plot we can clearly see decaying trend of non-seasonal lags after the lag 1, which is the biggest lag for the whole non-seasonal part and the subsequent lag experienced a sharp decline in ACF value hinting at possible MA(1) model (parameter $q = 1$). The same things can be said about seasonal lags in the ACF plot. There is a decaying trend of seasonal lags after the seasonal lag 1, which is the biggest seasonal lag for the whole seasonal part and the subsequent seasonal lag experienced a sharp decline in ACF value hinting at possible MA(1) model (parameter $Q = 1$). Now we move on to the PACF plot, which looks more complex than usual. The clearer trend is seen in seasonal lags of the PACF. The first seasonal lag is the biggest seasonal lag for the whole seasonal part and the subsequent seasonal lag experienced a sharp decline in PACF value hinting at possible AR(1) model (parameter $P = 1$). Here this trend is clearer than in the case

of seasonal lags in the ACF because after seasonal lag 4 values become small so that exceeding confidence intervals becomes nearly impossible. In the ACF plot we have seen some seasonal lags being slightly above confidence intervals even though their antecedent lags were way below confidence intervals. The proposal of non-seasonal AR models by solely looking at the PACF plot in this particular case seems non-trivial, hence it would be a good idea to propose several hypothetical non-seasonal AR models and later check them using several model validation techniques. The decaying trend of non-seasonal AR part is there, though it is slower and seems to start later comparing it to non-seasonal MA part. Therefore, one could argue that one possible non-seasonal AR model which is AR(2) (parameter $p = 2$) because lag 2 is followed a lag that has experienced a sharp decline in PACF value. However, lags 4,5, and 6 are definitely above confidence intervals. Here is where we can propose another non-seasonal AR mode which is AR(6) (parameter $p = 6$), because lag 6 is followed by lags that are below confidence intervals consequentially and the decline in their values is sharper. Regarding the values of parameters d and D , we don't need to propose them because we know the actual values which are $d = 1$ and $D = 1$. In the end, combining everything we have seen we can propose 9 SARIMA models in the order of increasing complexity, which are:

SARIMA Models

	Model	Description
1	SARIMA(0, 1, 1)(1, 1, 0) _{12}	Non-seasonal MA(1), Seasonal AR(1)
2	SARIMA(0, 1, 1)(0, 1, 1) _{12}	Non-seasonal MA(1), Seasonal MA(1)
3	SARIMA(0, 1, 1)(1, 1, 1) _{12}	Non-seasonal MA(1), Seasonal AR(1), Seasonal MA(1)
4	SARIMA(2, 1, 1)(1, 1, 0) _{12}	Non-seasonal MA(1), Non-seasonal AR(2), Seasonal AR(1)
5	SARIMA(2, 1, 1)(0, 1, 1) _{12}	Non-seasonal MA(1), Non-seasonal AR(2), Seasonal MA(1)
6	SARIMA(2, 1, 1)(1, 1, 1) _{12}	Non-seasonal MA(1), Non-seasonal AR(2), Seasonal AR(1), Seasonal MA(1)
7	SARIMA(6, 1, 1)(1, 1, 0) _{12}	Non-seasonal MA(1), Non-seasonal AR(6), Seasonal AR(1)
8	SARIMA(6, 1, 1)(0, 1, 1) _{12}	Non-seasonal MA(1), Non-seasonal AR(6), Seasonal MA(1)
9	SARIMA(6, 1, 1)(1, 1, 1) _{12}	Non-seasonal MA(1), Non-seasonal AR(6), Seasonal AR(1), Seasonal MA(1)

Below is the same data given in a more tabular manner.

```
##      p i q P D Q      AIC      BIC
## 10 6 1 1 0 1 1 470.7858 504.4439
## 6  2 1 1 0 1 1 471.1388 489.8378
## 12 6 1 1 1 1 1 472.6071 510.0050
## 8  2 1 1 1 1 1 473.0751 495.5138
## 2  0 1 1 0 1 1 473.8614 485.0808
## 4  0 1 1 1 1 1 475.6325 490.5917
## 7  2 1 1 1 1 0 517.4252 536.1242
## 11 6 1 1 1 1 0 517.7831 551.4413
## 3  0 1 1 1 1 0 521.1615 532.3809
## 5  2 1 1 0 1 0 580.9848 595.9440
```

```
## 9  6 1 1 0 1 0 587.1439 617.0623
## 1  0 1 1 0 1 0 594.9498 602.4294
```

From the output above we can state that our theory about possible non-seasonal AR(6) model might be correct. Some of the SARIMA models that have non-seasonal AR(6) part do very well in terms of AIC values, in fact the model SARIMA(6,1,1)(0,1,1)₁₂ has the lowest AIC value, however BIC penalizes complexity of the models differently than AIC does and in terms of BIC values these models are average at best. One might think that we should easily choose models that would have AIC values of 2-3 points higher but with BIC values of 10 points lower at worst than the earlier mentioned SARIMA models but AIC/BIC values is not the only indicator of selecting the most optimal time-series model. After experimenting with the proposed p, P, Q parameters for SARIMA models, we have decided to further validate 3 better than average models. These 3 models, `mod1`, `mod2`, and `mod3`, can be seen in the outputs above. Their value of p is different (6,2,0) but all other parameters are identical. We were thinking about inclusion of models that have seasonal AR part with the value of 1 (P = 1) but all such models always do worse in terms of AIC/BIC compared to exactly same models but with seasonal AR part with the value of 0 (P = 0). The first model `mod1` is the most complex in terms of parameters (the highest p parameter), it has the lowest AIC, and the highest BIC. The second model `mod2` is the 2nd ranked in terms of complexity of parameters (the second highest p parameter) and it has the second highest AIC and BIC. The third model `mod3` is the least complex in terms of parameters (the lowest p parameter), it has the highest AIC, and the lowest BIC.

2. Estimation

Based on the previous analysis, we decide to investigate three model and below are the coefficients and other information about the models.

2a.

```
##
## Call:
## arima(x = serie, order = c(6, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6          ma1          sma1
##      -0.0507  -0.0756  -0.0336  -0.1479  -0.0993  -0.1476  -0.4573  -0.7066
## s.e.   0.1745   0.0976   0.0706   0.0610   0.0646   0.0648   0.1705   0.0403
##
## sigma^2 estimated as 0.2441:  log likelihood = -226.39,  aic = 470.79

##
## Call:
## arima(x = serie, order = c(2, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12))
##
## Coefficients:
##          ar1          ar2          ma1          sma1
##      0.3098  0.1070  -0.8115  -0.7035
## s.e.  0.0912  0.0738   0.0696   0.0400
##
## sigma^2 estimated as 0.2508:  log likelihood = -230.57,  aic = 471.14

##
## Call:
```

```
## arima(x = serie, order = c(0, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12))
##
## Coefficients:
##          ma1      sma1
##       -0.5317 -0.7023
## s.e.    0.0599  0.0388
##
## sigma^2 estimated as 0.2565: log likelihood = -233.93, aic = 473.86
```

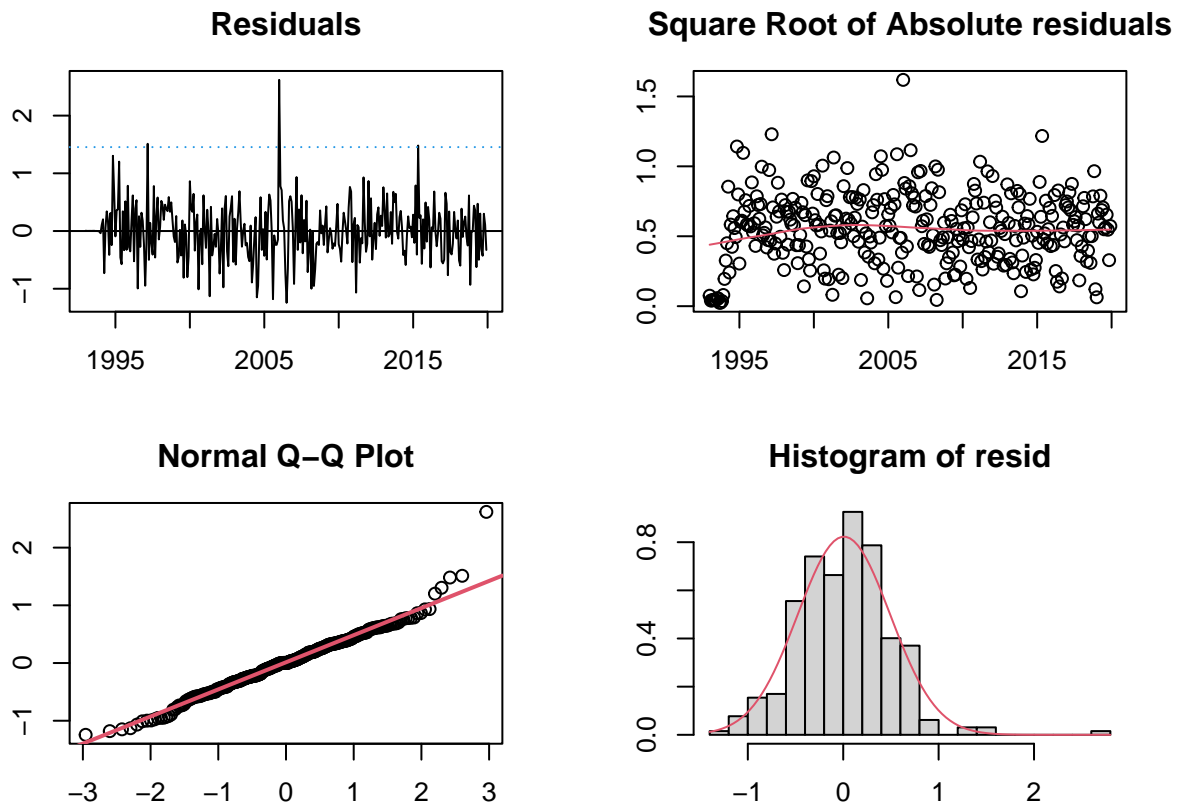
3. Validation

Now we will validate the proposed models, based on a thorough statistical analysis.


3a.

We will be taking a look at the outputs of the validation function for models `mod1`, `mod2`, and `mod3`. This function provides all the necessary things for the evaluation of a model which are complete analysis of residuals and roots of polynomials/ π and ψ weights to check the causality and/or invertibility.


```
validation(mod1)
```

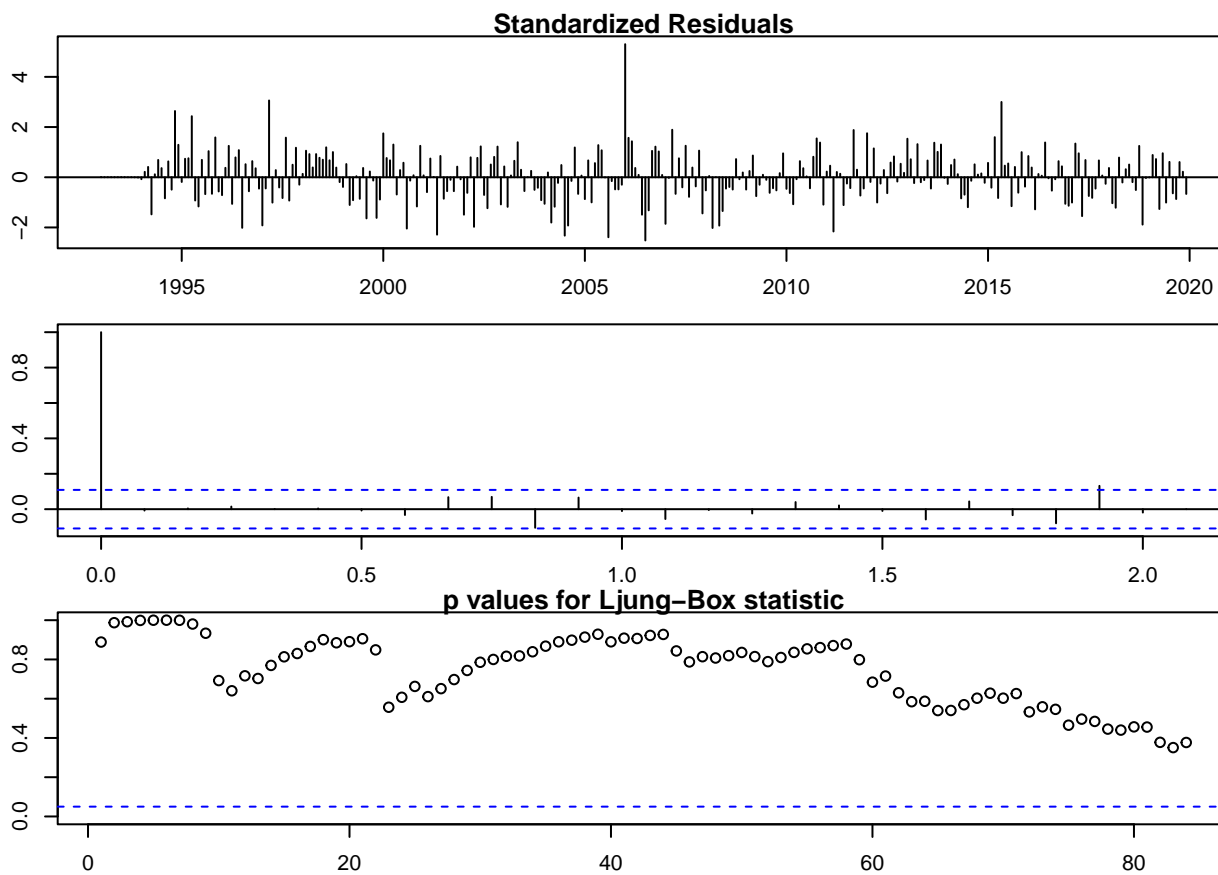


Series resid




Series resid





```
##
## -----
##
## Call:
## arima(x = serie, order = c(6, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12))
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ar5      ar6      ma1      sma1
##      -0.0507 -0.0756 -0.0336 -0.1479 -0.0993 -0.1476 -0.4573 -0.7066
## s.e.   0.1745   0.0976   0.0706   0.0610   0.0646   0.0648   0.1705   0.0403
##
## sigma^2 estimated as 0.2441:  log likelihood = -226.39,  aic = 470.79
##
## Modul of AR Characteristic polynomial Roots:  1.257526 1.421008 1.421008 1.257526 1.456458 1.456458
##
## Modul of MA Characteristic polynomial Roots:  1.029362 1.029362 1.029362 1.029362 1.029362 1.029362
```



```
##
## Psi-weights (MA(inf))
##
## -----
##      psi 1      psi 2      psi 3      psi 4      psi 5      psi 6
## -0.507935156 -0.049894727  0.007391469 -0.127446962 -0.016588236 -0.079591933
##      psi 7      psi 8      psi 9      psi 10     psi 11     psi 12
##  0.088411895  0.027576439  0.008601175  0.026744575 -0.005655552 -0.709748520
##      psi 13     psi 14     psi 15     psi 16     psi 17     psi 18
##  0.341543001  0.027683154 -0.006508707  0.088353367  0.015874257  0.059509588
##      psi 19     psi 20     psi 21     psi 22     psi 23     psi 24
## -0.059389174 -0.018531302 -0.006724334 -0.019685384  0.002659380  0.001431084
##
## Pi-weights (AR(inf))
##
## -----
##      pi 1      pi 2      pi 3      pi 4      pi 5      pi 6
## -0.507935156 -0.307892849 -0.174341419 -0.227608940 -0.203356548 -0.240617476
##      pi 7      pi 8      pi 9      pi 10     pi 11     pi 12
## -0.110024459 -0.050309652 -0.023004531 -0.010519024 -0.004809917 -0.708817453
##      pi 13     pi 14     pi 15     pi 16     pi 17     pi 18
## -0.359921847 -0.218022511 -0.123403072 -0.160928741 -0.143739378 -0.170044761
##      pi 19     pi 20     pi 21     pi 22     pi 23     pi 24
## -0.077754464 -0.035553913 -0.016257340 -0.007433812 -0.003399176 -0.500863409
##
## Descriptive Statistics for the Residuals
```

```

##
## -----
##          resid
## nobs      324.000000
## NAs        0.000000
## Minimum   -1.244061
## Maximum    2.619121
## 1. Quartile -0.304051
## 3. Quartile  0.329753
## Mean       0.007704
## Median     0.002521
## Sum        2.496165
## SE Mean    0.026928
## LCL Mean   -0.045272
## UCL Mean    0.060680
## Variance   0.234933
## Stdev      0.484699
## Skewness   0.422569
## Kurtosis   2.313684
##
## Normality Tests
##
## -----
## Shapiro-Wilk normality test
##
## data:  resid
## W = 0.97527, p-value = 2.291e-05
##
## Anderson-Darling normality test
##
## data:  resid
## A = 0.63104, p-value = 0.09927
##
## Jarque Bera Test
##
## data:  resid
## X-squared = 84.073, df = 2, p-value < 2.2e-16
##
## Homoscedasticity Test
##
## -----
## studentized Breusch-Pagan test
##
## data:  resid ~ I(obs - resid)
## BP = 0.00032599, df = 1, p-value = 0.9856
##
## Independence Tests
##

```



```
## -----
##
## Durbin-Watson test
##
## data: resid ~ I(1:length(resid))
## DW = 2.0149, p-value = 0.5313
## alternative hypothesis: true autocorrelation is greater than 0
##
##
## Ljung-Box test
##      lag.df    statistic    p.value
## [1,]      1  0.01964658 0.8885286
## [2,]      2  0.02575262 0.9872062
## [3,]      3  0.09764684 0.9921183
## [4,]      4  0.09821219 0.9988330
## [5,]     12  8.83949077 0.7165755
## [6,]     24 21.53189380 0.6071841
## [7,]     36 26.02163775 0.8898693
## [8,]     48 39.39546394 0.8072403
```


For the model SARIMA(6,1,1)(0,1,1)₁₂

- Residuals plot
 - The residuals are centered around zero with no evident trend
 - Occasional spikes above the confidence intervals that might be explain by potential outliers
- Square Root of Absolute Residuals
 - A reasonable consistent spread, which implies constant variance over time and homoscedasticity
- Normal Q-Q Plot
 - Small deviations from the diagonal line are seen at the ends of the line, but since the trend holds and deviations clearly can't be called heavy tails they might be caused by potential outliers
- Histogram of residuals
 - Some asymmetry is visible
 - A couple of observations above the bell-shaped curve are seen at the ends of it hinting at potential outliers
- ACF and PACF of Residuals
 - There are barely any lags that are above confidence intervals and if they are they are higher by negligible value indicating there are no significant autocorrelation/partial autocorrelations in residuals
- Standardized Residuals
 - The standardized residuals fluctuate around zero without evident patterns or trends, though some standardized residuals definitely look like spikes hinting at potential outliers
- p-values for Ljung-Box Test
 - None of the p-values for the Ljung-Box test are below 0.05, indicating no significant autocorrelation in the residuals.


Based on the SARIMA(6,1,1)(0,1,1)₁₂ model modules of AR/MA characteristic polynomial roots we can say that the model is both causal and invertible as all modules of polynomial roots are higher than 1.

```
validation(mod2)
```


Residuals




Square Root of Absolute residuals




Normal Q-Q Plot




Histogram of resid




Series resid




Series resid





```
##
## -----
##
## Call:
## arima(x = serie, order = c(2, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12))
##
## Coefficients:
##          ar1      ar2      ma1      sma1
##      0.3098  0.1070 -0.8115 -0.7035
## s.e.  0.0912  0.0738  0.0696  0.0400
##
## sigma^2 estimated as 0.2508:  log likelihood = -230.57,  aic = 471.14
##
## Modul of AR Characteristic polynomial Roots:  1.935003 4.83079
##
## Modul of MA Characteristic polynomial Roots:  1.029741 1.029741 1.029741 1.029741 1.029741 1.029741
```



```
##
## Psi-weights (MA(inf))
##
## -----
##      psi 1      psi 2      psi 3      psi 4      psi 5
## -0.5017111600 -0.0484454876 -0.0686806734 -0.0264592282 -0.0155442118
##      psi 6      psi 7      psi 8      psi 9      psi 10
## -0.0076460279 -0.0040315706 -0.0020669062 -0.0010716011 -0.0005530873
##      psi 11     psi 12     psi 13     psi 14     psi 15
## -0.0002859799 -0.7036440546  0.3528755715  0.0340417549  0.0482962030
##      psi 16     psi 17     psi 18     psi 19     psi 20
##  0.0186034284  0.0109298481  0.0053761371  0.0028347401  0.0014533090
##      psi 21     psi 22     psi 23     psi 24
##  0.0007534788  0.0003888941  0.0002010821  0.0001038968
##
## Pi-weights (AR(inf))
##
## -----
##      pi 1      pi 2      pi 3      pi 4      pi 5      pi 6
## -0.50171116 -0.30015958 -0.24357972 -0.19766513 -0.16040540 -0.13016911
##      pi 7      pi 8      pi 9      pi 10     pi 11     pi 12
## -0.10563233 -0.08572072 -0.06956243 -0.05644996 -0.04580919 -0.74067048
##      pi 13     pi 14     pi 15     pi 16     pi 17     pi 18
## -0.38311882 -0.23564160 -0.19122334 -0.15517788 -0.12592697 -0.10218983
##      pi 19     pi 20     pi 21     pi 22     pi 23     pi 24
## -0.08292713 -0.06729543 -0.05461029 -0.04431629 -0.03596270 -0.52409079
```

```

##
## Descriptive Statistics for the Residuals
##
## -----
##               resid
## nobs          324.000000
## NAs            0.000000
## Minimum       -1.576139
## Maximum        2.679299
## 1. Quartile   -0.294924
## 3. Quartile    0.322825
## Mean          0.007232
## Median         0.008385
## Sum            2.343328
## SE Mean       0.027295
## LCL Mean      -0.046466
## UCL Mean       0.060931
## Variance       0.241388
## Stdev          0.491313
## Skewness       0.315620
## Kurtosis       2.492236
##
## Normality Tests
##
## -----
## Shapiro-Wilk normality test
##
## data:  resid
## W = 0.97518, p-value = 2.203e-05
##
## Anderson-Darling normality test
##
## data:  resid
## A = 0.77975, p-value = 0.04255
##
## Jarque Bera Test
##
## data:  resid
## X-squared = 91.589, df = 2, p-value < 2.2e-16
##
## Homoscedasticity Test
##
## -----
## studentized Breusch-Pagan test
##
## data:  resid ~ I(obs - resid)
## BP = 0.14472, df = 1, p-value = 0.7036
##
##

```

```
## Independence Tests
##
## -----
##
## Durbin-Watson test
##
## data: resid ~ I(1:length(resid))
## DW = 2.0095, p-value = 0.5118
## alternative hypothesis: true autocorrelation is greater than 0
##
##
## Ljung-Box test
##      lag.df      statistic    p.value
## [1,]      1  0.007915379 0.9291070
## [2,]      2  0.012102978 0.9939668
## [3,]      3  1.400260430 0.7054737
## [4,]      4  2.795663177 0.5925816
## [5,]     12 17.690224319 0.1254252
## [6,]     24 31.895893189 0.1295954
## [7,]     36 36.764672815 0.4332839
## [8,]     48 50.200280730 0.3862616
```

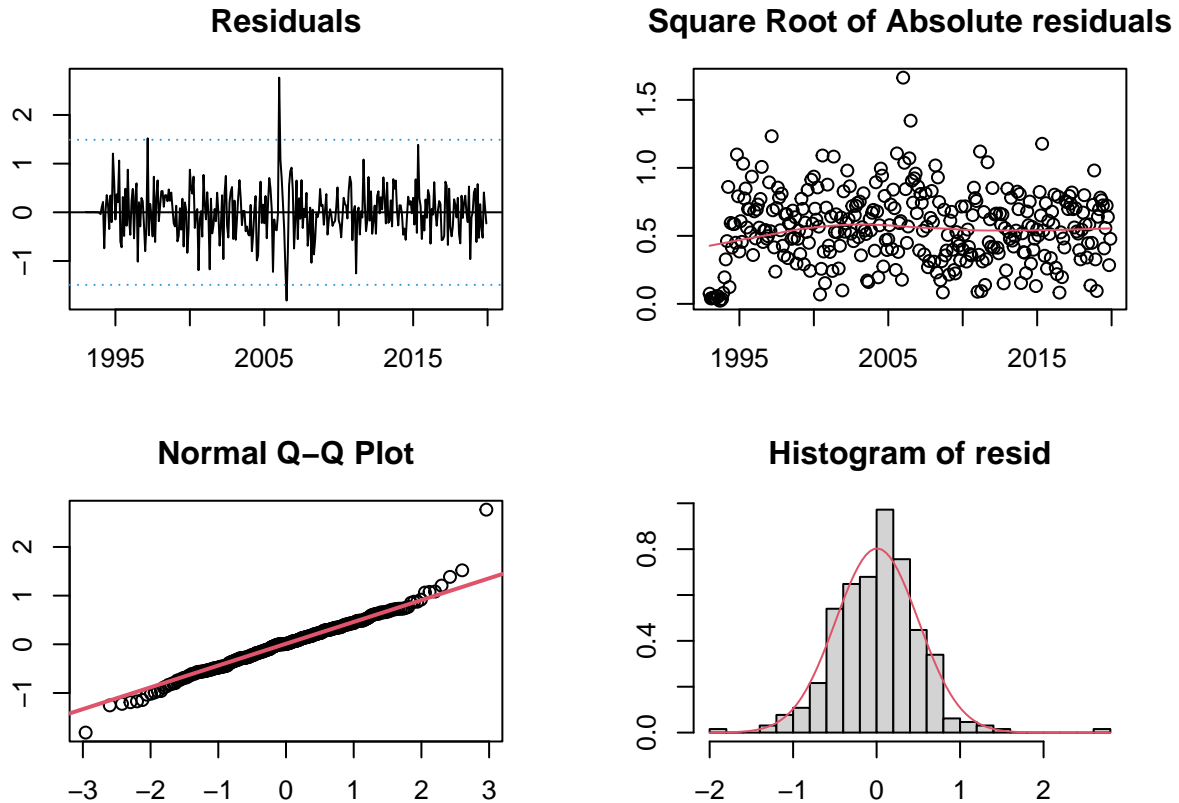
For the model SARIMA(2,1,1)(0,1,1)₁₂

- Residuals plot
 - The residuals are centered around zero with no evident trend
 - Occasional spikes above the confidence intervals that might be explain by potential outliers
- Square Root of Absolute Residuals
 - A reasonable consistent spread, which implies constant variance over time and homoscedasticity
- Normal Q-Q Plot
 - Small deviations from the diagonal line are seen at the ends of the line, but since the trend holds and deviations clearly can't be called heavy tails they might be caused by potential outliers
 - Points at both ends of the lined distributed worse compared to SARIMA(6,1,1)(0,1,1)₁₂ model
 - Even though they are most likely called by potential outliers it means that skewness and kurtosis values might be different for this model compared to SARIMA(6,1,1)(0,1,1)₁₂ model
- Histogram of residuals
 - Some asymmetry is visible
 - A couple of observations above the bell-shaped curve are seen at the ends of it hinting at potential outliers
- ACF and PACF of Residuals
 - There are barely any lags that are above confidence intervals and if they are they are higher by negligible value indicating there are no significant autocorrelation/partial autocorrelations in residuals
- Standardized Residuals
 - The standardized residuals fluctuate around zero without evident patterns or trends, though some standardized residuals definitely look like spikes hinting at potential outliers
- p-values for Ljung-Box Test


- None of the p-values for the Ljung-Box test are below 0.05 and the lowest p-value is around 0.0525, indicating no significant autocorrelation in the residuals.
- Compared to SARIMA(6,1,1)(0,1,1)12 model all p-values are significantly lower across the board, though they are still above the confidence interval

Based on the SARIMA(2,1,1)(0,1,1)12 model modules of AR/MA characteristic polynomial roots we can say that the model is both causal and invertible as all modules of polynomial roots are higher than 1.


```
validation(mod3)
```




Series resid

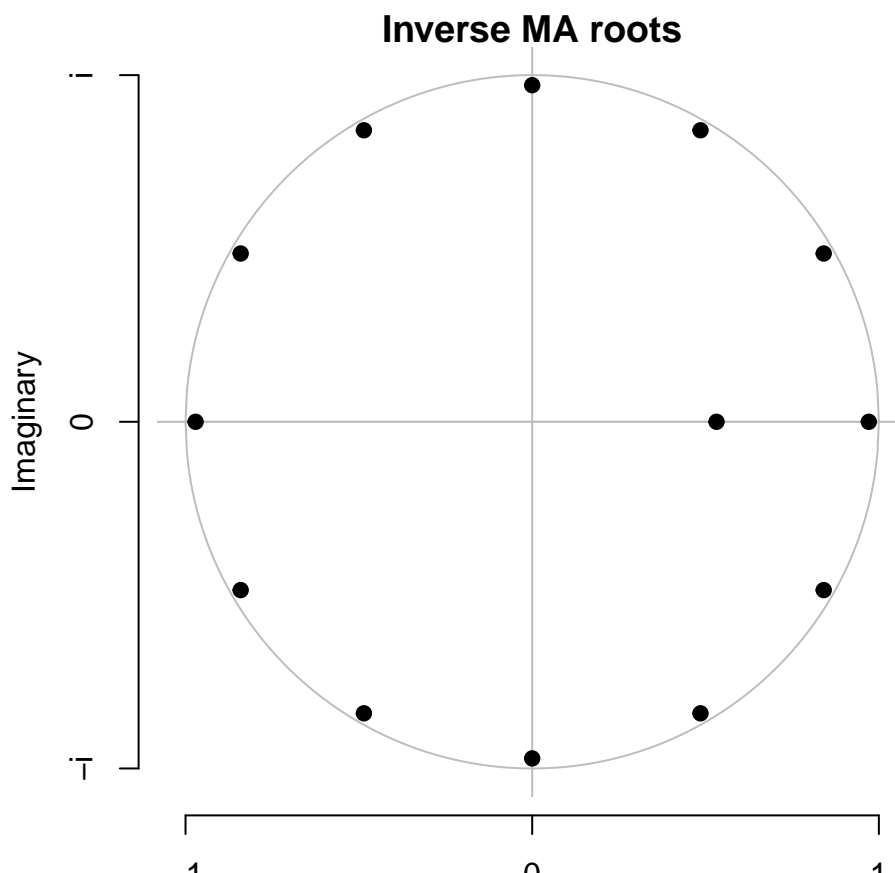


Series resid





```
##
## -----
##
## Call:
## arima(x = serie, order = c(0, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12))
##
## Coefficients:
##          ma1      sma1
##      -0.5317  -0.7023
## s.e.   0.0599   0.0388
##
## sigma^2 estimated as 0.2565:  log likelihood = -233.93,  aic = 473.86
##
## Modul of AR Characteristic polynomial Roots:
##
## Modul of MA Characteristic polynomial Roots:  1.029882 1.029882 1.029882 1.029882 1.029882 1.029882
```



```
##
## Psi-weights (MA(inf))
##
## -----
##      psi 1      psi 2      psi 3      psi 4      psi 5      psi 6      psi 7
## -0.5316889  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
##      psi 8      psi 9      psi 10     psi 11     psi 12     psi 13     psi 14
##  0.0000000  0.0000000  0.0000000  0.0000000 -0.7023469  0.3734301  0.0000000
##      psi 15     psi 16     psi 17     psi 18     psi 19     psi 20     psi 21
##  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
##      psi 22     psi 23     psi 24
##  0.0000000  0.0000000  0.0000000
##
## Pi-weights (AR(inf))
##
## -----
##      pi 1      pi 2      pi 3      pi 4      pi 5
## -0.5316889043 -0.2826930910 -0.1503047798 -0.0799153837 -0.0424901228
##      pi 6      pi 7      pi 8      pi 9      pi 10
## -0.0225915268 -0.0120116642 -0.0063864686 -0.0033956145 -0.0018054105
##      pi 11     pi 12     pi 13     pi 14     pi 15
## -0.0009599167 -0.7028572828 -0.3737014186 -0.1986928978 -0.1056428091
##      pi 16     pi 17     pi 18     pi 19     pi 20
## -0.0561691094 -0.0298644923 -0.0158786192 -0.0084424856 -0.0044887759
##      pi 21     pi 22     pi 23     pi 24
## -0.0023866324 -0.0012689459 -0.0006746845 -0.4936498983
```

```

##
## Descriptive Statistics for the Residuals
##
## -----
##          resid
## nobs      324.000000
## NAs        0.000000
## Minimum   -1.813913
## Maximum    2.766604
## 1. Quartile -0.290170
## 3. Quartile  0.314306
## Mean       0.005771
## Median     0.008353
## Sum        1.869924
## SE Mean    0.027605
## LCL Mean   -0.048537
## UCL Mean    0.060080
## Variance   0.246899
## Stdev      0.496889
## Skewness   0.335869
## Kurtosis   2.914760
##
## Normality Tests
##
## -----
## Shapiro-Wilk normality test
##
## data:  resid
## W = 0.97334, p-value = 1.054e-05
##
## Anderson-Darling normality test
##
## data:  resid
## A = 0.80512, p-value = 0.03683
##
## Jarque Bera Test
##
## data:  resid
## X-squared = 123.75, df = 2, p-value < 2.2e-16
##
## Homoscedasticity Test
##
## -----
## studentized Breusch-Pagan test
##
## data:  resid ~ I(obs - resid)
## BP = 0.087161, df = 1, p-value = 0.7678
##
##

```

```
## Independence Tests
##
## -----
##
## Durbin-Watson test
##
## data: resid ~ I(1:length(resid))
## DW = 1.9072, p-value = 0.1856
## alternative hypothesis: true autocorrelation is greater than 0
##
##
## Ljung-Box test
##      lag.df  statistic    p.value
## [1,]      1  0.7024856 0.401949748
## [2,]      2  0.8220396 0.662973808
## [3,]      3  0.9231266 0.819843553
## [4,]      4  7.3925133 0.116543473
## [5,]     12 27.7554998 0.006006002
## [6,]     24 41.3896224 0.015089534
## [7,]     36 47.2073674 0.100083041
## [8,]     48 58.7986444 0.136583780
```

For the model SARIMA(0,1,1)(0,1,1)₁₂

- Residuals plot
 - The residuals are centered around zero with no evident trend
 - Occasional spikes above the confidence intervals that might be explain by potential outliers
- Square Root of Absolute Residuals
 - A reasonable consistent spread, which implies constant variance over time and homoscedasticity
- Normal Q-Q Plot
 - Small deviations from the diagonal line are seen at the ends of the line, but since the trend holds and deviations clearly can't be called heavy tails they might be caused by potential outliers
 - SARIMA(0,1,1)(0,1,1)₁₂ model has the best Normal Q-Q plot
- Histogram of residuals
 - Some asymmetry is visible
 - A couple of observations above the bell-shaped curve are seen at the ends of it hinting at potential outliers
- ACF and PACF of Residuals
 - There are barely any lags that are above confidence intervals and if they are they are higher by negligible value indicating there are no significant autocorrelation/partial autocorrelations in residuals
- Standardized Residuals
 - The standardized residuals fluctuate around zero without evident patterns or trends, though some standardized residuals definitely look like spikes hinting at potential outliers
- p-values for Ljung-Box Test
 - Around 50% of the p-values for the Ljung-Box test are below 0.05 and the lowest p-value is around close to 0 (somewhere around 0.01-0.02)

- For the most part p-values struggle to be above 0.05 for higher lags indicating significant serial autocorrelation

Based on the SARIMA(0,1,1)(0,1,1)₁₂ model modules of AR/MA characteristic polynomial roots we can say that the model is both causal and invertible as all modules of polynomial roots are higher than 1.

Now we are going to check the stability of the proposed models and evaluate their capability of prediction, reserving the last 12 observations.

```
# Correct split for training and testing
ultim <- c(2018, 12) # Train until the end of 2018
serie_train <- window(serie, end = ultim) # Training data (1993 to 2018)
serie_test <- window(serie, start = ultim + c(0, 1)) # Test data (2019)

# Fit models to training data
(mod1B <- arima(serie_train, order = c(6, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12)))

##
## Call:
## arima(x = serie_train, order = c(6, 1, 1), seasonal = list(order = c(0, 1, 1),
##      period = 12))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6          ma1          sma1
##      -0.0239  -0.0688  -0.0227  -0.1430  -0.1075  -0.1269  -0.4723  -0.7019
## s.e.   0.1839   0.1003   0.0730   0.0625   0.0664   0.0684   0.1793   0.0419
##
## sigma^2 estimated as 0.2479:  log likelihood = -220.09,  aic = 458.18

(mod2B <- arima(serie_train, order = c(2, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12)))

##
## Call:
## arima(x = serie_train, order = c(2, 1, 1), seasonal = list(order = c(0, 1, 1),
##      period = 12))
##
## Coefficients:
##          ar1          ar2          ma1          sma1
##      0.3230   0.1007  -0.8111  -0.6998
## s.e.  0.0937   0.0752   0.0720   0.0414
##
## sigma^2 estimated as 0.2541:  log likelihood = -223.75,  aic = 457.51

(mod3B <- arima(serie_train, order = c(0, 1, 1), seasonal = list(order = c(0, 1, 1), period = 12)))

##
## Call:
## arima(x = serie_train, order = c(0, 1, 1), seasonal = list(order = c(0, 1, 1),
##      period = 12))
##
## Coefficients:
##          ma1          sma1
```

```
##          -0.5182  -0.6979
## s.e.    0.0618   0.0403
##
## sigma^2 estimated as 0.2601:  log likelihood = -227.11,  aic = 460.21
```


```
# Predict for the next 12 months (2019)
pred1 <- predict(mod1B, n.ahead = 12)
pred2 <- predict(mod2B, n.ahead = 12)
pred3 <- predict(mod3B, n.ahead = 12)

# Convert predictions and standard errors into time series
pr1 <- ts(pred1$pred, start = ultim + c(0, 1), freq = 12)
se1 <- ts(pred1$se, start = ultim + c(0, 1), freq = 12)
pr2 <- ts(pred2$pred, start = ultim + c(0, 1), freq = 12)
se2 <- ts(pred2$se, start = ultim + c(0, 1), freq = 12)
pr3 <- ts(pred3$pred, start = ultim + c(0, 1), freq = 12)
se3 <- ts(pred3$se, start = ultim + c(0, 1), freq = 12)

# Compute confidence intervals for predictions
tl1 <- pr1 - 1.96 * se1
tu1 <- pr1 + 1.96 * se1
tl2 <- pr2 - 1.96 * se2
tu2 <- pr2 + 1.96 * se2
tl3 <- pr3 - 1.96 * se3
tu3 <- pr3 + 1.96 * se3


# Visualize predictions for Model 1 with Actual Data
ts.plot(serie, tl1, tu1, pr1, lty = c(1, 2, 2, 1), col = c(1, 4, 4, 2),
        xlim = c(2017, 2020), type = "o",
        main = paste("Model ARIMA(6,1,1)(0,1,1)12"))
abline(v = (2017:2020), lty = 3, col = 4)
```

Model ARIMA(6,1,1)(0,1,1)12




```
# Visualize predictions for Model 2
ts.plot(serie, tl2, tu2, pr2, lty = c(1, 2, 2, 1), col = c(1, 4, 4, 2),
        xlim = c(2017, 2020), type = "o",
        main = paste("Model ARIMA(2,1,1)(0,1,1)12"))
abline(v = (2017:2020), lty = 3, col = 4)
```


Model ARIMA(2,1,1)(0,1,1)12



```
# Visualize predictions for Model 3
ts.plot(serie, tl3, tu3, pr3, lty = c(1, 2, 2, 1), col = c(1, 4, 4, 2),
        xlim = c(2017, 2020), type = "o",
        main = paste("Model ARIMA(0,1,1)(0,1,1)12"))
abline(v = (2017:2020), lty = 3, col = 4)
```

Model ARIMA(0,1,1)(0,1,1)12



```
# Evaluate Prediction Errors (2019 Observations)
obs <- window(serie, start = ultim + c(0, 1)) # Observed data for 2019
error1 <- obs - window(pr1, start = ultim + c(0, 1))
error2 <- obs - window(pr2, start = ultim + c(0, 1))
error3 <- obs - window(pr3, start = ultim + c(0, 1))

# Compute Metrics for Each Model
metrics <- data.frame(
  Model = c("ARIMA(6,1,1)(0,1,1)", "ARIMA(2,1,1)(0,1,1)", "ARIMA(0,1,1)(0,1,1)"),
  RMSE = c(sqrt(mean(error1^2, na.rm = TRUE)),
            sqrt(mean(error2^2, na.rm = TRUE)),
            sqrt(mean(error3^2, na.rm = TRUE))),
  MAE = c(mean(abs(error1), na.rm = TRUE),
            mean(abs(error2), na.rm = TRUE),
            mean(abs(error3), na.rm = TRUE)),
  RMSPE = c(sqrt(mean((error1 / obs)^2, na.rm = TRUE)),
              sqrt(mean((error2 / obs)^2, na.rm = TRUE)),
              sqrt(mean((error3 / obs)^2, na.rm = TRUE))),
  MAPE = c(mean(abs(error1 / obs), na.rm = TRUE),
            mean(abs(error2 / obs), na.rm = TRUE),
            mean(abs(error3 / obs), na.rm = TRUE))
)

# Display Metrics
metrics
```

##	Model	RMSE	MAE	RMSPE	MAPE
## 1	ARIMA(6,1,1)(0,1,1)	0.3382936	0.2877056	0.02888241	0.02451385
## 2	ARIMA(2,1,1)(0,1,1)	0.3492316	0.2881464	0.02958092	0.02448518
## 3	ARIMA(0,1,1)(0,1,1)	0.3678022	0.2910628	0.03068842	0.02447672